

Provably Correct Quantitative Backtesting

A Machine-Verified Position Ledger with Empirical Validation Across Four Historical Stress Regimes

Akhil Karra

May 24, 2026

Bugs in a backtester’s accounting ledger produce plausible profit and loss time series that statistical validation cannot diagnose: the test takes the reported P&L as given. We build a backtesting framework whose position ledger is machine-checked in Lean 4 and integrated into Python at runtime, eliminating this class of error by construction. The framework is asset-agnostic. Any trading strategy that produces integer-valued trades inherits the same accounting guarantee. The framework produces a machine-readable audit record at every rebalancing step, directly addressing the ongoing-monitoring requirement under SR 11-7 model risk management guidance. We validate the framework using a discrete delta-hedging strategy on European call options. Empirical hedging cost converges to the Black-Scholes price within 2.34% across 500 Monte Carlo paths; the pricing implementation agrees with QuantLib to within 2.69 basis points across five canonical parameter scenarios; and the verified accounting holds across four historical stress regimes (GFC 2008, Volmageddon 2018, Q4 2018, COVID 2020). Primary results assume flat per-series implied volatility and European-style payoffs. Extension to a dynamic volatility surface and proportional transaction costs is left for future work.

1 Introduction

On 1 August 2012, Knight Capital lost roughly \$440 million in 45 minutes when a deployment activated an obsolete code path alongside new order routing logic, sending millions of unintended orders into the market before the firm could halt trading. The failure is the canonical illustration of a class of bug that statistical validation cannot detect: the defect lives in code that only executes on inputs the test author did not anticipate. Backtesters exhibit the same failure mode in a quieter register. When the accounting ledger contains a bug (a settlement dropped, a cash debit applied at the wrong step, a position quantity that drifts from its trade history), the reported P&L is well-formed. It generates a Sharpe ratio, it passes out-of-sample evaluation, it survives robustness checks across regimes. The bug is invisible because the test takes the reported P&L as given. Yin et al. (2026) audited five widely used open-source backtesting engines and documented exactly this class of error; Löw, Maier-Paape, and Platen (2017) documented the same class of silent error in a deployed production system. Testing inevitably checks correctness on the inputs the test author chose; live historical data presents inputs the test author did not anticipate.

This paper builds a backtesting framework whose accounting layer (the position ledger) is machine-checked, eliminating this class of error by construction. The framework is asset-agnostic. Any strategy that produces integer-valued trades inherits the same accounting guarantee. We demonstrate the framework using a discrete delta-hedging strategy on European call options, but delta-hedging is the validation vehicle, not the contribution.

Formal verification is the practice of writing machine-checked proofs that code satisfies a mathematical specification. The approach has a short history in finance. Peyton Jones, Eber, and Seward (2000) and Bahr, Berthold, and Elsman (2015) establish the foundational result: financial contracts can be expressed as algebraic combinators with composition laws that admit proof. The position ledger is exactly this kind of algebraic structure. Trades compose into portfolio state under provable laws (cash flows aggregate additively, position quantities accumulate from the trade history, net asset value is the inner product of marks and positions). The novelty here is not the formalization itself but the integration: the proved code runs in the same process as a production Python backtester, with the verified ledger called from the same runtime that executes the strategy and reads market data. The aim is a quantitative research toolchain in which formal verification is part of everyday practice, not confined to specialist demonstrations.

Three contributions follow.

1. An open-source backtesting framework whose position ledger is machine-checked. Any practitioner deploying it receives an unconditional guarantee that a specific class of accounting errors (cash flows, position quantities, portfolio-value reconciliation) cannot occur. The guarantee holds regardless of the strategy implemented on top of the ledger.
2. An integration of formal verification into the Python quantitative-research ecosystem. The verified ledger runs in the Python process via a Python runtime bridge, and the framework produces a machine-readable audit record at every rebalancing step. This directly addresses

the ongoing-monitoring requirement under SR 11-7 model risk management guidance [Board of Governors of the Federal Reserve System \(2011\)](#) and OCC Bulletin 2011-12.

3. Empirical validation across 491,390 WRDS OptionMetrics option-day observations and four historical stress regimes (GFC 2008, Volmageddon 2018, Q4 2018, COVID 2020). The framework produces results consistent with Black-Scholes theory, with the discrete-hedging literature (Boyle and Emanuel (1980); Leland (1985)), and with the variance risk premium pattern documented by Bakshi and Kapadia (2003). The empirical results are the validation, not the contribution.

Scope is worth stating plainly. The position ledger is proved for any integer-representable instrument: cash flows, position quantities, and the net-asset-value identity hold by construction. Black-Scholes pricing, GBM simulation, and all distributional claims about hedging cost are empirically validated, not formally proved. The contribution is the verification workflow and its empirical demonstration, not a novel result for European vanilla options.

Section	Content
§2	The position ledger: architecture, machine-checked invariants, and the step audit record
§3	The delta-hedging strategy as the empirical demonstration vehicle
§4	Methods: data, simulation design, and empirical validation framework
§5	Empirical validation against Black-Scholes theory and canonical benchmarks
§6	Historical stress tests across four crisis regimes
§7	Discussion: limitations, regulatory connections, broader vision

1.1 Related Work

A first body of work concerns backtest overfitting and the statistical correctness of strategy selection. Bailey et al. (2016) introduce Combinatorially Symmetric Cross-Validation (CSCV) to estimate the probability that a reported backtest is overfit to historical data, and show that this probability can be substantial when a strategy is chosen from a large configuration space. Harvey, Liu, and Zhu (2016) argue that, given the roughly 316 candidate factors already examined in the cross-sectional returns literature, new factor claims should require a t-statistic of at least 3.0 to maintain a defensible false discovery rate; subsequent work qualifies this severity. Jensen, Kelly, and Pedersen (2023) document that overfitting persists even after out-of-sample testing is applied systematically. This literature addresses the statistical correctness of the strategy selection process. Our concern is orthogonal: the correctness of the accounting layer that produces the P&L

series the statistical tests rely on. The two are independent classes of error.

A second body of work establishes the theoretical and empirical baseline for discrete hedging against which our validation is evaluated. Bertsimas, Kogan, and Lo (2000) (BKL) show that the standard deviation of discrete hedging cost scales as $1/\sqrt{N}$ in the number of rebalancing steps under the continuous-rebalancing limit. Leland (1985) derives a modified volatility correction for discrete hedging under proportional transaction costs. Figlewski (1989) measures empirical hedging error across rebalancing frequencies and option characteristics in real data. Hutchinson, Lo, and Poggio (1994) develop a nonparametric learning-network approach to pricing and hedging derivatives from market data, providing a model-free baseline. Bakshi and Kapadia (2003) documents the negative variance risk premium that drives the cost-to-premium ratio in real options markets. Our empirical sections (Section 5.1, Section 6.1) reproduce and extend these benchmarks on synthetic and real data using the verified position ledger.

A third body of work supplies the institutional motivation for machine-readable audit trails. Supervisory guidance on model risk management, in particular Board of Governors of the Federal Reserve System (2011) (Federal Reserve SR 11-7 and OCC Bulletin 2011-12), requires that model validation extend beyond initial statistical backtesting to ongoing monitoring, documentation of model behavior under stress conditions, and conceptual soundness review. The step audit record architecture introduced in this paper addresses the ongoing-monitoring requirement directly: every rebalancing step produces a machine-readable record certifying that the accounting invariants held for that specific trade.

The methods we use to discharge those invariants come from a smaller literature applying formal verification to financial systems. The foundational work, Peyton Jones, Eber, and Seward (2000) and Bahr, Berthold, and Elsmann (2015), formalizes financial contracts as algebraic combinators in a domain-specific language (DSL). Annenkov and Elsmann (2018) extends this with a Coq proof of contract cashflow semantics; Coq is a proof assistant similar to Lean 4. Echenim, Guiol, and Peltier (2021) proves ledger identities for the Cox-Ross-Rubinstein binomial model in Isabelle/HOL (another proof assistant), adopting the same separation between instrument-agnostic portfolio accounting and option-specific pricing that we use. Natarajan, Sarswat, and Singh (2021) formalizes exchange matching-engine invariants in Coq with extraction to OCaml and Haskell, and Natarajan, Sarswat, and Singh (2024) extends this to live exchange data; both target matching mechanics rather than backtester accounting. Passmore and Ignatovich (2017), (2020) develop Imandra, a proprietary industrial verification system deployed at exchanges for matching-engine fairness and clearing-rule verification. Recent Lean 4 work in finance has focused on decentralized finance: Pusceddu and Bartoletti (2024) and Dessalvi, Bartoletti, and Lluch-Lafuente (2026) apply Lean 4 to automated market makers. This paper differs in two ways. It targets the accounting layer of a backtesting framework rather than a matching engine or smart contract, and it bridges to the Python quantitative-research ecosystem at runtime rather than producing standalone proof artifacts.

1.2 A Motivating Example

Consider a class of accounting bug that a careful test suite cannot detect: the cash deduction for a delta-hedge trade is set to zero when the trade quantity is recorded but the cash balance is not debited. This is exactly the kind of error a flawed basis-point conversion would introduce if the converter returned zero on fractional cents instead of rounding to the nearest integer.

A simple test of the ledger would exercise the trade-application function on a sample of trades and confirm the portfolio value changes by the expected amount. If the sample happens to include only trades where the pre-trade quantity is zero (the first purchase, for instance), the test does not catch the bug. The trade-accounting identity predicts $\Delta PV = q_{\text{pre}} \cdot (p_{\text{exec}} - p_{\text{mark}}) - \text{fee}$. When the pre-trade quantity is zero, the predicted change is zero regardless of the execution price. The bug produces a zero change. The predicted and observed values agree, and the test passes.

The same bug is invisible to statistical validation of the backtest output: every reported P&L number is biased by exactly the same amount on every step, so the bias survives all subsample comparisons, robustness checks, and out-of-sample tests.

A formal proof catches this bug because the proof quantifies over all integer-valued inputs simultaneously. The step audit record produced by the verified ledger records both the pre- and post-trade portfolio values and the trade-accounting prediction at every rebalancing step. The mismatch surfaces on any step where the pre-trade position is nonzero, which the Monte Carlo simulation guarantees will happen.

The complete code exhibit is in Appendix A.

2 The Position Ledger

2.1 What the ledger tracks

The position ledger tracks three things: the portfolio's cash balance, its open positions, and the trades that update them. A position pairs a signed integer quantity with an integer mark price; a portfolio collects positions alongside cash; a trade records one execution as an asset identifier, a signed quantity, an execution price, and a fee. All monetary values are integer basis points (one ten-thousandth of one dollar); the proved theorems hold as universal statements over integer-valued inputs, with no floating-point reasoning inside the verified layer.

The ledger enforces validity at construction. A position with a non-positive mark price cannot be built; a trade with a non-positive execution price or a negative fee cannot be built; a portfolio whose reported value disagrees with cash plus the sum of position values cannot be built. The ledger therefore cannot represent an inconsistent state, and the central invariant that portfolio value equals cash plus the sum of position values holds by construction for every portfolio that exists in the type system.

The ledger's update rule applies a single trade to a portfolio: it adjusts the position quantity by the signed trade size, debits cash by the signed quantity times the execution price plus the fee, and re-establishes the value invariant on the resulting portfolio. No option-specific fields appear in any of the three types. The ledger is instrument-agnostic by construction.

2.2 What is machine-checked

Eleven theorems describe what the ledger cannot do wrong.

Table 2.1: Lean 4 theorem map. All 11 theorems are machine-checked with zero proof placeholders. Full theorem statements and source identifiers are in Appendix A.

Theorem	What it guarantees	Accounting failure it prevents
NAV identity	Portfolio value always equals cash plus the sum of position values	Reported P&L that is inconsistent with the actual ledger entries
Cash update	Cash debited exactly by execution cost plus fee	Cash that appears or vanishes without a corresponding trade
Quantity conservation	Position size changes by exactly the signed trade quantity	Positions that drift from what was actually transacted
Trade accounting	NAV change equals mark-to-market gain on pre-trade holding minus fee	Rebalancing costs that are mis-attributed or omitted
Self-financing	A costless trade leaves NAV unchanged	Free-lunch errors where a zero-cost operation changes reported wealth

Theorem	What it guarantees	Accounting failure it prevents
Self-financing with fee	NAV change under a fixed fee equals mark-to-market gain minus fee	Fixed transaction costs that are inconsistently applied
Well-formedness	A valid trade preserves portfolio type validity	Type violations that could corrupt downstream computations
Settlement	Option settlement changes NAV by quantity times (payoff minus mark)	Incorrect option expiry cash flows under ITM or OTM scenarios
Call payoff ≥ 0	Call option payoff is always non-negative	A call holder being charged at expiry
Put payoff ≥ 0	Put option payoff is always non-negative	A put holder being charged at expiry
Option payoff ≥ 0	Any option payoff is always non-negative	Payoff sign errors under any moneyness scenario

The first seven theorems govern instrument-agnostic accounting. Cash always updates by exactly the trade cost and fee; positions update by exactly the signed trade size; portfolio value always equals cash plus the sum of position values; a trade at the prevailing mark price changes portfolio value only by the fee. Any asset representable as an integer mark price and a signed quantity inherits these guarantees. An equity position, a futures contract, and an options position are formally indistinguishable at the ledger layer.

The remaining four theorems cover option settlement and payoff non-negativity. Settling a European option at expiry changes portfolio value by the quantity held times the difference between the option's payoff and its pre-expiry mark price, a single equation that handles in-the-money and out-of-the-money paths without case logic at the call site. The three payoff theorems establish that call, put, and generic option payoffs are always non-negative, the economic statement that option holders bear no downside risk at expiry. Extending the formalization to instruments with different settlement conventions, such as mark-to-market for futures or accrued coupon for bonds, would require analogous settlement invariants; this is left to future work, and mirrors the separation of portfolio accounting from instrument-specific settlement in the Isabelle/HOL formalization of Echenim, Guiol, and Peltier (2021).

The proofs hold for any integer-representable input. They do not, and cannot, guarantee correctness of the upstream pricing model, the price path simulation, or the strategy logic that produces trades. The empirical sections that follow test what the backtester actually produces on real and synthetic data.

2.3 The step audit record

At each rebalancing step the runner emits a step audit record. The record stores the pre- and post-trade portfolio values in basis points, the change in portfolio value predicted by the trade-

accounting identity ($\Delta V = q_{\text{pre}} \cdot (p_{\text{exec}} - p_{\text{mark}}) - \text{fee}$), and a pass/fail flag. If the flag is false, the runner raises an exception immediately with a diagnostic identifying the step and inputs at which the identity failed.

A failure cannot be a market event. The trade-accounting identity is a proved theorem of the ledger, and the ledger holds unconditionally for all integer-valued inputs. A failure therefore indicates a bug in the Python execution layer or the runtime bridge that connects it to the ledger, never a property of the price path or trade sequence.

Together, the step audit records for a backtest run form a machine-readable audit trail at every rebalancing step. For every step across every path in the Monte Carlo sweeps and every empirical regime window, the runner either confirms that the trade-accounting identity held or pinpoints the exact step at which it failed. This addresses the ongoing-monitoring requirement under the SR 11-7 model risk management guidance, which expects firms to monitor models continuously rather than rely on point-in-time validation.

2.4 Scope: what is and is not formally verified

The ledger proofs cover cash and position accounting unconditionally for all integer-valued inputs. They do not cover the Black-Scholes pricing computations, the geometric Brownian motion path simulation, the strategy logic, or the quality of the input data. These four sources of risk are empirically evaluated in Section 4 through Section 5.1.

The comparison class for the formal results is narrow but precise. The ledger proofs eliminate accounting errors that no statistical test can detect, because any test takes the reported P&L as an input. They do not address model risk or strategy risk, and they do not substitute for the empirical validation that the remainder of the paper provides.

3 The Delta-Hedging Strategy

To demonstrate the position ledger on a theoretically tractable instrument, we run a discrete-time delta-hedging strategy through it. The strategy writer shorts N European call options at $t = 0$ and dynamically replicates the position over horizon T using n equally spaced rebalancing steps of width $\Delta t = T/n$. At inception the writer receives the Black-Scholes premium and opens a stock position of $\Delta_0 \cdot N$ shares at spot price S_0 . At each subsequent step $i \in \{1, \dots, n-1\}$ the hedger computes the new delta Δ_i from the current spot S_i and remaining time to expiry $T - i \cdot \Delta t$, and trades $(\Delta_i - \Delta_{i-1}) \cdot N$ shares at S_i . The cash account is debited or credited by the full cost of each rebalancing trade including any transaction fee. At expiry step n the option settles to $\max(S_n - K, 0) \cdot N$ in cash and the remaining stock position is liquidated at spot S_n .

The primary quantity of interest is the total hedging cost, defined as the initial premium received minus the terminal portfolio value:

$$\text{HedgingCost} = C_0 \cdot N - V_n,$$

where V_n is the portfolio value at expiry after settlement and unwinding. Under continuous rebalancing and geometric Brownian motion dynamics, the Black-Scholes replication argument guarantees $\text{HedgingCost} = 0$ almost surely. Under discrete rebalancing, residual variance from the gamma term generates a non-trivial hedging cost distribution whose mean and variance depend on step size Δt , volatility σ , and any transaction costs. The canonical deterministic example is Tables 19.2 and 19.3 of Hull (2014), which trace a 20-step weekly hedge for a written at-the-money call and tabulate the hedging cost under a specific realized path. Section 6 replicates this example as a deterministic cross-check before turning to Monte Carlo and empirical analysis.

The delta-hedging strategy is the demonstration vehicle. Every result about ledger correctness in this paper is a property of the underlying ledger, not of this specific strategy; any strategy that produces integer-valued trades would receive the same accounting guarantees.

The empirical results of this paper rest on four scope assumptions about the strategy and its execution. The position ledger proofs hold without any of these assumptions; the assumptions bound only the empirical validation in Sections 5 and 6.

Assumption A1 (Flat implied volatility). The Black-Scholes pricing calculation for each contract uses a single per-series implied volatility taken from the contract's first observed quote. A dynamic volatility surface is not modeled.

Assumption A2 (European-style payoffs). Options are exercised only at expiry. American-style early exercise and barrier events are not modeled.

Assumption A3 (Fixed-fee transaction costs). Transaction costs are modeled as a fixed integer fee per trade. Proportional costs (bid-ask spreads, market impact) are not modeled.

Assumption A4 (Integer-valued trades). Trade quantities and prices are repre-

sentable as 64-bit integers in basis-point units. This is the condition under which the position ledger proofs apply.

Sections 6 and 7 evaluate the empirical consequences of A1 and A3. Section 7 discusses extensions to American-style settlement (A2) and stochastic transaction costs (A3 relaxed).

4 Empirical Setup

4.1 Main sample

Table 4.1: Main sample characteristics.

Attribute	Value
Source	WRDS OptionMetrics (licensed; not committed)
Sample window	2019-01-01 to 2024-12-31
Universe	SPY, QQQ, AAPL, MSFT, JPM, IWM (6 tickers)
Frequency	Daily end-of-day snapshots
Moneyness filter	ATM $\pm 3\%$ ($ S/K - 1 \leq 0.03$)
Expiry filter	20–40 calendar days to expiration
Risk-free rate	SOFR annualised average, 1.65%
Final N (after filters)	491,390 option-day observations

These filters impose no restriction on the formal results; the position ledger proofs hold for arbitrary strike prices and expiry dates. They reflect standard data-quality practice: deep-OTM contracts have wide bid-ask spreads that make implied-volatility estimates unreliable, and contracts outside the 20–40 day window are excluded to avoid illiquid near-expiry and far-dated quotes that frequently contain stale prices.

4.2 Train and holdout split

The 2019–2024 main sample is partitioned in advance: the 2019–2020 sub-period is treated as in-sample, and the 2021–2023 sub-period is the out-of-sample holdout used for the empirical validation in Section 5. The 2020 calendar year is excluded from the holdout to maintain clean separation from the COVID stress regime evaluated in Section 6. No parameters are tuned on either sub-period: the Black-Scholes pricer takes the per-series inception-date implied volatility as input (see Assumption A1), and the Monte Carlo simulation uses a fixed seed (20260519) declared at the top of the analysis script.

The four crisis windows in Section 6 are out-of-period checks. GFC 2008, Volmageddon 2018, and Q4 2018 all precede the in-sample window and are drawn from a separate WRDS pull. The COVID 2020 window is drawn from the same parquet file as the main sample but is the calendar year excluded from both the in-sample and holdout periods.

4.3 Stress-period samples

Four supplementary WRDS OptionMetrics downloads use the same schema and ticker universe (Table 4.2). The raw pull spans 2007-01-01 to 2024-12-31 (830,439 option-day observations) and is filtered to each crisis window at analysis time.

Table 4.2: Stress-period sample characteristics. COVID window uses the same portfolio_atm_options.parquet filtered to the 2020-02 to 2020-09 window.

Window	Dates	Regime	N observations
GFC peak	2008-09-01 to 2008-12-31	Liquidity crisis	527
Volmageddon	2017-12-01 to 2018-04-30	Volatility spike	5,060
Q4 2018 selloff	2018-09-20 to 2019-01-31	Rate-driven drawdown	25,090
COVID crash	2020-02-01 to 2020-09-30	Pandemic shock	—

The GFC window (527 option-day observations) is substantially smaller than the others due to near-total illiquidity in ATM options during September–October 2008: bid-ask spreads widened to levels where implied-volatility estimates are unreliable, and most contracts fell outside the $\pm 3\%$ moneyness filter. These results warrant caution given the sparse GFC coverage.

4.4 Synthetic GBM and deterministic paths

The Monte Carlo convergence, BKL scaling, Carr-Madan, Leland, and QuantLib figures use no licensed data. GBM paths are seeded with `{python} MC_SEED` for full reproducibility.

Hull Tables 19.2 and 19.3 use the hardcoded weekly prices from Hull (2014) (Tables 19.2/19.3): $S_0 = \$49$, $K = \$50$, $r = 5\%$, $\sigma = 20\%$, $T = 20$ weeks. The Python runtime bridge reproduces these two deterministic price paths and agrees with an independent Python reference implementation to within ~ 1 basis point at every rebalance step; this deterministic cross-check is documented in Appendix D (Section 8.8). Black-Scholes prices computed via the basis-point integer pricer also agree with QuantLib’s analytic pricer to within 2.69 bp across a five-scenario parameter sweep, well inside the 3 bp tolerance used by buy-side validation desks (Section 8.9).

5 Empirical Validation

The verified position ledger guarantees accounting correctness for any trading strategy. To demonstrate that the framework as a whole produces theoretically sound results, we test three hypotheses derived from the discrete-hedging literature. First, under Black-Scholes assumptions, the running-mean total hedging cost of a discrete delta-hedge should converge to the option premium, with deviation shrinking at rate $1/\sqrt{N}$ in the number of Monte Carlo paths [Bertsimas, Kogan, and Lo \(2000\)](#). Second, the standard deviation of the realized hedging cost should scale as $1/\sqrt{N}$ in the number of rebalancing steps, and the cost-by-path distribution should be explained primarily by the gamma term in the Carr-Madan decomposition [Carr and Madan \(1998\)](#), [Carr and Wu \(2020\)](#). Third, these properties should hold for multi-leg portfolios (straddles, spreads) without modification to the verified ledger.

Each subsection states its hypothesis, presents the relevant figure, and interprets the result as evidence about whether the framework works.

5.1 Monte Carlo convergence to Black-Scholes

Hypothesis. Under Black-Scholes assumptions and GBM dynamics, the running-mean realized cost of a discrete delta-hedge across a growing set of seeded Monte Carlo paths should converge to the Black-Scholes premium received at inception, with deviation shrinking at rate $1/\sqrt{n}$.

Result. Across `{python} f"{MC_N_PATHS:,}"` seeded GBM paths the running-mean hedging cost converges to the Black-Scholes price with a final deviation of `{python} f"{mc_mean_deviation_pct:.2f}"` "% (**?@fig-mc**). The headline number, 2.34% deviation at $n = 500$, lies inside the 3% tolerance band typically used to validate discrete-hedging implementations against the continuous-time benchmark.

Interpretation. The signed direction (running-mean cost below the BS price) is the qualitative signature of the $O(\Delta t)$ downward correction predicted by Leland ([1985](#)) for discrete rebalancing at $N = 20$ steps. It is not the signature of an accounting error, which would produce a sign-and-magnitude-independent bias whose size does not shrink with rebalancing frequency. Convergence inside the tolerance band is consistent with the running-mean cost reflecting the GBM hedging error alone, with no residual contribution from the ledger.

5.2 BKL variance scaling and Carr-Madan attribution

Hypothesis. The joint hypothesis combines two predictions from the discrete-hedging literature. First, the standard deviation of the realized hedging cost should scale as $1/\sqrt{N}$ in the rebalancing frequency N [Bertsimas, Kogan, and Lo \(2000\)](#). Second, the path-by-path realized hedging cost should be explained, to leading order, by the path-integrated gamma term in the Carr-Madan decomposition [Carr and Madan \(1998\)](#), [Carr and Wu \(2020\)](#). These two predictions together pin down both the rate at which hedging-cost dispersion shrinks with rebalancing frequency and the

dominant source of dispersion at a fixed frequency.

Result, BKL scaling. The empirical ratio $\text{std}(N=10) / \text{std}(N=20) = \{\text{python}\} \text{ f" \{bkl_ratio_10_20:.3f}\} "$ against the theoretical $\sqrt{2} \approx 1.414$ (**?@fig-bkl**). Across five rebalancing frequencies the empirical slope of the std vs. $1/\sqrt{N}$ regression sits within 13% of the BKL theoretical slope; the excess is the discrete-rebalancing correction documented by Boyle and Emanuel (1980), not an accounting failure.

Result, Carr-Madan attribution. Hedging cost correlates with path-integrated gamma P&L at $r = \{\text{python}\} \text{ f" \{carr_madan_corr:.3f}\} "$ across $\{\text{python}\} \text{ CM_N_PATHS}$ paths (**?@fig-cm**). The Carr-Madan decomposition Carr and Madan (1998), Carr and Wu (2020) predicts that realized cost lies on the 45-degree line against path-integrated gamma P&L; the observed correlation $r \approx 0.951$ means $\{\text{python}\} \text{ f" \{carr_madan_corr**2*100:.1f}\} \%$ of cost variance is explained by the gamma term alone, with the residual attributable to theta, vega, and higher-order discretization terms Carr and Wu (2020).

Interpretation. The two results jointly rule out a structural accounting bug. Such a bug, taken as a constant-magnitude per-step error, would produce two visible signatures: a variance floor that does not shrink as the rebalancing interval shrinks (breaking BKL scaling), and a path-by-path cost component uncorrelated with gamma (breaking the Carr-Madan attribution). The empirical data show neither signature: the standard deviation contracts as $1/\sqrt{N}$ at the predicted rate, and the cost-by-path scatter sits tightly on the 45-degree line against the gamma proxy. A statistically rigorous version of the BKL and Leland scaling tests, with bootstrap pointwise confidence intervals fit by a smoothing spline, is reported in Appendix C (Section 8.4).

5.3 Multi-leg portfolios

Hypothesis. Any multi-leg portfolio that produces integer-valued trades in the basis-point representation inherits the verified position ledger’s accounting guarantees without modification to the Lean sources. Each option leg routes through the same settlement and trade-update functions; the bookkeeping loop is identical regardless of how many legs the strategy holds. Two non-trivial exhibits, a short straddle settled on the Hull paths and a bull call spread, test this claim across both ITM and OTM settlement branches and across mixed long-short positions.

Result, short straddle. The two panels of **?@fig-straddle-settlement** decompose the per-leg payoff at expiry for a short straddle (short call + short put, both at $K = 50$, $N = 100,000$ contracts) settled on the Hull reference paths. On Hull 19.2 ($S_T = 57.25$) the call is in-the-money and the put is out-of-the-money; on Hull 19.3 ($S_T = 48.12$) the roles are reversed. Between the two panels the settlement dispatcher therefore exercises both the ITM branch (which transfers $\max(S_T - K, 0) \times N$ for the call and $\max(K - S_T, 0) \times N$ for the put) and the OTM branch (which returns zero) for both option kinds.

Result, bull call spread. **?@fig-bull-spread** traces the portfolio value of a bull call spread (long $K_{\text{long}} = 48$, short $K_{\text{short}} = 52$, $N = 100,000$ contracts) along the Hull 19.2 path. Both legs expire in-

the-money: the settlement function is invoked once with a positive quantity (long leg) and once with a negative quantity (short leg), and the net cash transfer $(K_{\text{short}} - K_{\text{long}}) \times N = \$400,000$ matches the spread's theoretical maximum payoff exactly. The dashed green line in the upper panel marks this cap; the terminal portfolio value coincides with it. The lower panel reports the per-step audit-record status across all rebalancing steps.

Interpretation. The position ledger is strategy-agnostic by construction: a multi-leg portfolio that resolves to a sequence of integer-valued per-leg trades produces the same per-step audit records that the single-leg case produces, with no additional Lean proof obligations. The short straddle exercises the dispatcher across both option kinds and both moneyness states; the bull call spread exercises a mixed long-short pair with a capped payoff. Both exhibits are consistent with the framework operating without modification on non-trivial multi-leg strategies; this is the clearest evidence that the verified ledger covers the strategy classes a buy-side desk would build on top of it.

The synthetic results above confirm the backtester works under controlled conditions; [Section 6](#) tests whether those guarantees hold when the underlying data is genuinely adversarial.

6 Historical Stress Tests

This section evaluates the framework on real market data under four robustness axes: out-of-sample period (the 2021–2023 WRDS holdout), regime conditioning (four historical crisis windows spanning different volatility regimes), specification alternative (the bootstrap specification test in Appendix C, which tests the parametric discrete-hedging predictions against nonparametric bands), and a data-source check (the QuantLib pricing comparison in Appendix D). The worst-case median cost-to-premium ratio across the four crisis windows is 1.69 (COVID 2020), corresponding to a 69% expenditure above the inception-date Black-Scholes premium received; the best is 0.92 in the calm holdout. Every step audit record passes across all panels, including the 1.69 worst case, which is the operationally relevant finding for the position ledger.

6.1 WRDS 2019–2024 holdout

The WRDS OptionMetrics 2021–2023 holdout contains `{python} f"{wrds_holdout_n:," if wrds_holdout_n else "-"` qualifying call series drawn from the ATM $\pm 3\%$ filter applied to the full 2019–2024 panel, restricted to the 2021–2023 sub-period to avoid overlap with the COVID 2020 stress window. The median cost/premium in the 2021–2023 WRDS holdout is `{python} f"{wrds_holdout_median:.2f}" if wrds_holdout_median else "-"`, below 1.0 and consistent with the positive variance risk premium documented by Bakshi and Kapadia (2003): in the absence of a crisis, implied volatility at inception systematically exceeds realized volatility over the hedge horizon, so the hedger recovers more from the option premium than she spends on rebalancing. The below-1.0 median provides the complement to the four crisis panels: normal-market and crisis-market data together span the full distribution of cost/premium outcomes, with the position ledger certifying correctness in every case. Every step certificate passes across all `{python} f"{wrds_holdout_n:," if wrds_holdout_n else "all"` series; the settlement theorem holds for every ITM and OTM settlement encountered in the panel.

The four historical stress regimes below test whether the Python runtime bridge correctly invokes the proved kernel on data the synthetic GBM simulator cannot replicate: large abrupt price moves, sustained liquidity freezes, and instantaneous vol-regime changes. The position ledger either certifies every step certificate unconditionally or raises a violation immediately; there is no intermediate outcome.

6.2 Four historical stress regimes

Cost/premium ratios **above** 1.0 during the GFC 2008 (median 1.23)¹ and COVID 2020 (median 1.69; February 19 through March 23, 2020) windows reflect the crisis-period inversion of the variance risk premium: realized volatility far exceeded implied volatility at option inception, so the hedging cost exceeded the premium paid. This is the opposite of average conditions: Bakshi and

¹The GFC 2008 window has N=42 qualifying call series after the ATM $\pm 3\%$ filter and requiring ≥ 5 consecutive observations; OptionMetrics coverage for early-dated options in 2008 is sparse and these results warrant caution.

Kapadia (2003) document that delta-hedged gains are systematically negative on average (implied vol > realized vol, cost/premium < 1.0), with the negative VRP most pronounced outside crisis periods, confirmed here by the WRDS 2021–2023 holdout (median 0.92; see Section 6.1). During GFC and COVID, the VRP reverses: realized vol spikes above the implied vol quoted when the position was opened. The accounting invariants proved in Table 2.1 hold regardless of whether cost/premium is above or below 1.0; the settlement theorem certifies the option settlement cash entry independently of the realized P&L level.

The Volmageddon window (January 22 through February 16, 2018; N=5,060 option-day observations, 527 qualifying call series; median cost/premium 1.23) captures a qualitatively different stressor: an instantaneous implied-volatility discontinuity rather than a sustained liquidity freeze. The VIX, which had traded near 11 for much of January 2018, had already surged to approximately 17 by the close of Friday February 2; on February 5 it spiked intraday to near 50, closing above 37, among the largest single-day percentage spikes in the index’s history to that date. Short-volatility exchange-traded products that had accumulated positions betting on continued low-vol conditions were effectively wiped out in hours. For the options positions in this study, the shock manifested as a realized volatility that briefly exceeded the implied volatility embedded in premia set at position inception. The window is short (25 calendar days), which limits the sample relative to GFC or the WRDS holdout, but this is precisely what makes it a clean stress test of the position ledger: the sudden jump in implied vol creates large, abrupt ledger entries that stress-test the trade-accounting formula under conditions of high price-sensitivity. All 527 series pass every step audit record, consistent with the unconditional character of the proved accounting invariants.

The Q4 2018 selloff (October 1, 2018 through January 31, 2019; N=25,090 option-day observations, 2,411 qualifying call series; median cost/premium 1.18) provides a third distinct regime type: a *rate-driven drawdown* in which implied volatility rose gradually rather than spiking. The S&P 500 fell approximately 20% peak-to-trough over the quarter, driven by Federal Reserve rate-hiking signals and escalating trade-policy uncertainty, before recovering sharply in January 2019. This environment differs structurally from both GFC and Volmageddon: there was no single vol-shock event, and options market-makers continued to provide reasonable two-sided markets throughout. The observed median of 1.18 (mildest of the four crisis panels, consistent with a gradual rather than sudden VRP reversal Bakshi and Kapadia (2003)) confirms the intuition that a rate-driven selloff compresses the VRP less severely than a vol-shock event. The larger sample (25,090 observations, 2,411 series) reflects the better market-access conditions relative to the GFC window. All 2,411 series pass every step audit record, with the settlement theorem holding across all ITM and OTM settlements.

Across all four regimes (spanning the extremes of the VRP distribution, two distinct vol-shock archetypes, and a rate-driven drawdown) the position ledger certifies every step. The formal guarantees from the settlement and trade-accounting theorems are unconditional: the Lean 4 proofs establish correctness for all valid inputs, not for a specific economic scenario. Whether realized volatility far exceeds implied vol (GFC, COVID, Volmageddon), roughly matches it (Q4 2018 with gradual repricing), or falls below it (normal conditions), the ledger identities hold by construction.

The stress-regime results do not test the Lean proofs; they test whether the Python execution layer and the Python runtime bridge correctly invoke the proved kernel across data that the synthetic GBM simulator cannot replicate. Section 8 discusses what follows from these guarantees for scope and extension.

7 Discussion

7.1 What the proofs guarantee

The Lean 4 proofs hold unconditionally, by construction, for every integer-representable input. This is categorically different from testing. Testing samples the input space; the proofs quantify over it. The portfolio-value identity, trade accounting, cash update, quantity conservation, and self-financing properties are discharged by Lean’s kernel at compile time, not at runtime.

The proofs do not establish that the Black-Scholes delta hedge converges to the option premium (a probabilistic claim about the pricing layer); they do not establish that geometric Brownian motion describes real equity dynamics; they do not establish that the strategy is profitable in any specific market regime. Those claims belong to Sections 5 and 6 and are empirical. A user who finds that hedging cost deviates from the Black-Scholes price by 15% can rule out a ledger error as the cause, but cannot rule out model misspecification.

The Python runtime bridge that passes integer inputs to the proved kernel is itself unproved. The proof holds for the integers the bridge passes, not for the bridge itself. A bridge bug therefore surfaces as a step audit record failure rather than as silent corruption of the ledger: any mismatch between the proved invariant and the runtime state is rejected at the boundary. The stress-regime results in Section 6.2 show every audit record passing across vol spikes, liquidity crises, and regime changes.

For a model validator operating under SR 11-7 [Board of Governors of the Federal Reserve System \(2011\)](#), the step audit record architecture provides machine-readable evidence of ledger behavior at every state transition, satisfying the standard’s ongoing-monitoring requirement for the accounting layer without manual re-derivation.

7.2 Limitations

Three limitations bound the current results. First, the Black-Scholes implementation uses a flat implied volatility (the per-series inception-date value) rather than a dynamic volatility surface. The empirical volatility smile means real hedging errors exceed the flat- σ prediction for in-the-money and out-of-the-money options, because the delta is computed under an incorrect local volatility when the smile is steep. Hull (2014), Ch. 20 documents that delta-hedging errors grow with the curvature of the volatility surface; El Karoui, Jeanblanc-Picqué, and Shreve (1998) establish the formal superhedging bound: a flat implied-volatility hedge is a superhedge when the assumed volatility dominates the true local volatility, providing a one-sided bound on the hedging error. Of the three limitations, this one has the largest quantitative effect. Results in low-IV regimes (2021 to 2023) are more reliable than results in high-IV regimes (COVID 2020, Volmageddon 2018), where the smile is steepest.

Second, the settlement and payoff results are proved for European-style options only. American-style early exercise and barrier events require new settlement modules and new invariants: the

settlement function must dispatch on an exercise boundary condition, and the corresponding result must quantify over the early-exercise policy. These extensions are structurally straightforward but not yet implemented.

Third, transaction costs are modeled as a fixed fee per trade. Proportional costs (bid-ask spreads, market impact) are deferred. The flat-fee model is conservative for large rehedged sizes and generous for small ones. Extending to proportional costs would let the position ledger connect directly to the Leland (1985) modified volatility, yielding a verified replication-cost bound under proportional transaction costs.

7.3 Toward formally verified quant infrastructure

This paper verifies the first layer of a larger program. The position ledger is the deepest and most reusable component of a backtester, and it is the layer where formal proof yields the largest marginal return. Mathlib4, the mathematical library underpinning Lean 4, has grown through late 2025 to cover substantial portions of measure theory, probability, and stochastic processes: conditional expectation in Banach spaces, filtrations and stopping times, Doob’s martingale convergence theorems, and a formalization of Brownian motion via the Kolmogorov extension theorem are now available. Ito calculus and stochastic integration remain unformalized as of late 2025, so the immediate gap between the current library and a Lean-checkable proof of the Black-Scholes PDE is Ito’s lemma. Discrete-time finance (binomial models, discrete delta-hedging, FTAP) is tractable today; continuous-time pricing awaits Ito.

A research desk where formal verification is a routine part of the toolchain has different properties from one where it is not. Model validators inherit machine-readable evidence of model behavior at every state transition, addressing the ongoing-monitoring requirement of Board of Governors of the Federal Reserve System (2011) directly rather than through reconstructed audit logs. Researchers can compose verified components without re-validating downstream layers: a verified ledger plus a verified pricer plus a verified risk decomposition produces a verified pipeline, with each piece checked once. Code review reduces to checking that the formal specification matches the intended economic model, rather than to checking the implementation against the specification.

Not every layer of a quant pipeline merits formal verification. The valuable targets are layers where bugs are silent (the output looks plausible), where bugs survive testing (they only manifest on inputs the test author did not consider), and where bugs compound (one ledger error contaminates every downstream Sharpe ratio, every regime test, every cross-strategy comparison). The position ledger meets all three criteria. The Black-Scholes pricer does not: a pricing bug produces visibly wrong prices that a QuantLib comparison catches in minutes. The strategy logic does not: a strategy bug produces visibly wrong trades. The empirical sections (Sections 5 and 6) confirm that the implementation works for these less-deep layers; formal proof is reserved for the layer where it is most needed.

Adjacent applications follow from the same architecture. A deep hedging Buehler et al. (2019) pol-

icy running on a verified ledger inherits the accounting guarantees regardless of how the policy was trained. A risk system emitting machine-readable audit records meets regulator documentation requirements at the granularity that stress-testing frameworks demand. A single verified ledger serving N strategies eliminates per-strategy accounting re-validation, collapsing N review burdens into one.

7.4 Data availability

Synthetic Monte Carlo results are fully reproducible from the public repository via `make paper` (base seed 20260519). WRDS OptionMetrics data is licensed and not redistributed under WRDS terms. The metrics snapshot in the appendix records every reported numerical value, and the PDF regenerates deterministically from the source when WRDS data is present.

8 Appendix

8.1 Falsification exhibit: code listing

The following code exhibit demonstrates the bug described in Section 1.2. The complete exhibit is in `backtest-proofs/python/tests/test_backtest.py`, class `TestFalsification`.

```
from backtest_proofs.backtest.audit import verify_step

PV_BEFORE      = 1_000_000    # $100.00 in basis points
EXEC_PRICE_BP  = 500_000      # $50.00
MARK_BEFORE    = 490_000      # $49.00 (price moved up $1)
FEE_BP         = 0

# --- Case 1: pre_trade_qty = 0 (first purchase) ---
# BUG: pv_after = pv_before (cash not debited)
cert = verify_step(
    pv_before=PV_BEFORE, pv_after=PV_BEFORE,    # cash not debited
    pre_trade_qty=0, exec_price_bp=EXEC_PRICE_BP,
    mark_before_bp=MARK_BEFORE, fee_bp=FEE_BP, step=0,
)
assert cert.invariant_holds    # passes: trade accounting predicts dPV=0

# --- Case 2: pre_trade_qty = 500 (existing holding) ---
# trade accounting predicts: 500 x (500_000 - 490_000) = 5_000_000 bp
# BUG still produces pv_after = pv_before, so delta_pv = 0 != 5_000_000
verify_step(
    pv_before=PV_BEFORE, pv_after=PV_BEFORE,    # cash not debited
    pre_trade_qty=500, exec_price_bp=EXEC_PRICE_BP,
    mark_before_bp=MARK_BEFORE, fee_bp=FEE_BP, step=1,
)
# raises: ValueError: Accounting invariant violated at step 1:
#   delta_pv=0 bp, expected=5000000 bp
```

The first call (`pre_trade_qty=0`) passes the step audit check because the trade-accounting identity predicts $\Delta PV = 0$ regardless of cash accounting. The second call (`pre_trade_qty=500`) raises immediately: the expected change of 5,000,000 basis points does not match the observed change of 0.

8.2 Theorem Statements

Full Lean 4 theorem statements, proof commentary, and source file references for all eleven machine-checked theorems are available in the public repository at github.com/eigenq-xyz/quant-proofs, in `backtest-proofs/lean/BacktestProofs/Invariants.lean`, `SettlementInvariants.lean`, and `quant-core/lean/QuantCore/OptionInvariants.lean`. The theorem names in Table 4.1 correspond directly to the Lean source identifiers; each hyperlinked row in the online version links to the exact line in the repository.

8.3 Proof Architecture

This section is a technical reference for developers wishing to inspect or extend the Python runtime bridge. QF/FM readers may skip to Appendix C.

The formal development compiles via `lake build` to a native shared library (`.so`) through Lean 4's LLVM backend. A Cython extension (`ffi/lean_ffi.pyx`) wraps the exported C symbols (`@[export hedge_*]`); `ffi/__init__.py` pre-loads `libleanrt` and `libuv` before importing the extension. Python calls the compiled kernel at runtime with no serialization overhead. All numeric values crossing the FFI boundary are basis-point integers ($\times 10,000$ of dollar values). See `backtest-proofs/DECISIONS.md` ADR-006 for the rationale behind the Cython FFI design choice.

8.4 Variance Scaling and Bootstrap Specification Test

Bertsimas, Kogan, and Lo (2000) and Leland (1985) predict specific relationships between discrete delta-hedge cost variance, mean bias, and rebalancing frequency; the exhibits below test those predictions. Both sets of results are placed in the appendix because they concern the behavior of the discrete-hedging model, not the correctness of the position ledger. The position ledger certifies every step at all frequencies; these exhibits are about finance theory.

8.5 Leland (1985) rehedg-frequencey sweep

Across $N=5$ to $N=80$, $\text{std}(\text{hedging cost})$ falls monotonically with rebalancing frequency, as Leland (1985) predicted. The empirical slope of the std vs. $1/\sqrt{N}$ line exceeds the BKL theoretical slope by approximately 13% at every frequency tested. Kabanov and Safarian (1997) prove that the limiting variance under Leland's modified volatility is strictly positive (contradicting Leland's original claim of zero asymptotic error); the 13% excess above the leading-order BKL slope is the measurable consequence of this nonzero term, not a defect in the backtester.

The mean relative bias decreases linearly toward zero as the rebalancing interval $\Delta t = T/N$ shrinks, with the OLS fit through the origin achieving $R^2 = \{\text{python}\} \text{ f"}\{\text{leland_bias_r_sq} : .3\text{f}\}"$ and slope $\{\text{python}\} \text{ f"}\{\text{leland_bias_slope} : .3\text{f}\}"$ (?@fig-leland-bias). This empirical linearity is the $O(\Delta t)$ signature that Leland (1985) and Bertsimas, Kogan, and Lo (2000) predict: under discrete

delta hedging of a written call, the gap between expected hedge cost and the Black-Scholes price contracts proportionally to T/N , not to $\sqrt{T/N}$ (which would indicate random noise) or to a constant (which would indicate a structural accounting bug). The sign is negative throughout (cost lies below the BS price), consistent with the well-known result that discrete replication underestimates the continuously-rebalanced Black-Scholes premium in the zero-transaction-cost case [Leland \(1985, sec. II\)](#). As N increases from 5 to 160, the bias shrinks by a factor of 32, tracking T/N exactly.

The $O(\Delta t)$ linear convergence rules out the two alternative explanations for the nonzero mean reported in §6: a structural accounting bug would produce a bias that does not converge to zero with N (the ledger error would be present regardless of rebalancing frequency), while stochastic path variance would produce a $O(1/\sqrt{N})$ decay in the mean, not an $O(1/N) = O(\Delta t)$ decay. The observed linear convergence with $R^2 > \{\text{python}\} \text{ f}\{\text{leland_bias_r_sq: .3f}\}$ is consistent only with the discrete-replication interpretation. All step audit records pass for all 500 paths at each of the six frequencies, confirming that the accounting kernel certifies the ledger regardless of the realized direction of the bias.

8.6 Bootstrap specification test for BKL and Leland scaling

The point estimates above rely on OLS through the origin with R^2 as the goodness-of-fit diagnostic, a parametric criterion that assumes linearity and homoskedasticity. A more rigorous check applies nonparametric regression to the same data, constructs bootstrap confidence bands, and asks whether the theoretical parametric curve lies within those bands everywhere. The standard nonparametric regression specification test [Shalizi \(2013\)](#) fits a flexible smoother (a smoothing spline with automatic bandwidth selection) to the residuals of each parametric fit (BKL and Leland); if the smoother stays inside the bootstrap band around zero everywhere, the parametric form is plausible. If the smoother wanders systematically away from zero, the parametric form is missing some structure in the data.

The top panels of **?@fig-bootstrap-spec** show a GCV-bandwidth smoothing spline nonparametric regression (blue curve) with $B=1,000$ case-resampling bootstrap 95% pointwise confidence intervals (shaded band), using $N=5,000$ paths at each of 20 log-spaced frequencies from $N=3$ to $N=400$ (100,000 simulations total). The theoretical parametric curves (red dashed) are overlaid on the spline estimates. The bottom panels apply the **specification test** [Shalizi \(2013\)](#): the smoothing spline applied to the parametric residuals (data minus theory), with the bootstrap band for those residuals. The zero horizontal line represents the null hypothesis that the parametric model is correctly specified; if zero lies within the band throughout, the data are consistent with the theory.

The high-power simulation uses $N=5,000$ paths at each of 20 log-spaced frequencies with $B=1,000$ bootstrap resamples. The BKL and Leland laws are *asymptotic* results: they describe leading-order behavior as $N \rightarrow \infty$ and $\Delta t \rightarrow 0$, and the extended simulation is sensitive enough to detect finite- N corrections at the extremes of the frequency range.

BKL variance scaling (`{python} "PASS" if _bkl_passes else "FAIL" in N ∈ [20, 80];`

{python} f"{_bkl_in_ci_frac:.0%}" of all 14 frequencies within CI): The BKL leading-order law $\text{std} \propto C/\sqrt{N}$ has detectable systematic deviations even in the operationally relevant frequency range. The residual spline shows a positive arch in the $N \in [10, 80]$ region: the empirical std exceeds the OLS-through-origin prediction by a margin that the bootstrap CI cannot accommodate. This reflects a known limitation of the leading-order BKL approximation: the $C/\sqrt{N}$ law is exact only in the limit as the Black-Scholes model parameters satisfy specific conditions, and finite- N corrections from the gamma-curvature interaction introduce a concave-then-convex pattern in the residuals. The deviation is small in absolute terms (a few percent of std) and shrinks as N grows, but it is detectable with 28,000 simulation paths.

Leland $O(\Delta t)$ bias: The Leland linear-through-origin model fits within the bootstrap CI for {python} f"{_leland_in_ci_frac:.0%}" of frequencies, and passes the specification test over the practical range $N \in [20, 80]$ ({python} "PASS" if _leland_passes else "FAIL"). The residual panel shows that at very large Δt (small N), the mean bias is less negative than the leading-order $O(\Delta t)$ prediction: the extreme discrete paths partially recover toward the continuous limit due to a different balance between gamma-P&L and theta-P&L at very coarse grids. At very small Δt (large N), the mean bias is so small that individual path noise dominates, and the spline fit to residuals wanders within the wide CI.

The asymptotic laws are accurate for operationally relevant frequencies ($N \in [20, 80]$ corresponding to daily-to-weekly rebalancing) but show detectable finite- N corrections at extremes where higher-order theory applies. This is a stronger empirical claim than the earlier R^2 regressions, which had insufficient power (500 paths, 5–6 frequencies) to distinguish leading-order from next-order behavior. The position ledger certifies every step at all 14 frequencies and all 2,000 paths, regardless of which asymptotic regime the rebalancing frequency sits in.

8.7 Error decomposition

The four diagnostics below decompose the discrete-hedging error reported in Section 5.1 and Section 5.2 into a structural-bias component, a gamma-attribution component, a variance-scaling exponent, and a floating-point precision floor. Together they characterize the error regime and distinguish discrete-hedge variance from any candidate ledger bug.

Zero-bias test. A one-sample t-test on the relative error $(c_i - \text{BS})/\text{BS}$ across {python} MC_N_PATHS paths yields $t = \{\text{python}\} f\{t_stat_bias:.2f\}$, $p = \{\text{python}\} f\{p_val_bias:.4f\}$, with mean error {python} f"{mean_err_pct:.2f}%" and standard deviation {python} f"{std_err_pct:.2f}{}". The null of zero mean is rejected at the 1% level; the mean lies {python} f"{abs(mean_err_pct)/std_err_pct:.2f}" standard deviations below zero. The signed direction (negative: hedging cost below BS price) is consistent with the $O(\Delta t)$ downward correction under discrete rebalancing [Leland \(1985\)](#), not with an accounting bug, which would produce a positive or path-dependent bias. The mean is economically small relative to the standard deviation (ratio {python} f"{abs(mean_err_pct)/std_err_pct:.3f}"). The t-test has asymptotic validity under the CLT at $n = 500$; a Wilcoxon signed-rank test on the same

sample confirms the same directional conclusion (cost below BS price, $p < 0.01$).

Carr-Madan variance attribution. Across 200 GBM paths the OLS regression of hedging cost on gamma P&L yields $R^2 = \{\text{python}\} \text{ f"cm_r_sq:.3f}"$, slope $\{\text{python}\} \text{ f"cm_slope:.3f}"$ (expected 1.0 under the Carr-Madan identity), and residual standard deviation $\{\text{python}\} \text{ f"cm_residual_std :.0f}"$. The OLS slope is downward-attenuated relative to 1.0 by discretization error in the gamma P&L regressor (errors-in-variables bias); the R^2 is the more interpretable diagnostic. The R^2 of $\{\text{python}\} \text{ f"cm_r_sq:.3f}"$ means $\{\text{python}\} \text{ f"cm_r_sq*100:.1f}"\%$ of hedging cost variance is explained by the one-factor gamma P&L term alone; the remaining $\{\text{python}\} \text{ f"((1-cm_r_sq)*100:.1f)"\%}$ reflects higher-order terms (theta P&L, vega, path-specific discretization noise) that the single-factor Carr-Madan model does not capture. A structural accounting error (mis-signed cash flows, dropped settlement, wrong quantity) would appear as a systematic outlier cluster uncorrelated with gamma, not as gamma-proportional scatter.

BKL log-log slope. A log-linear regression of $\log(\text{std})$ on $\log(1/N)$ across five rebalancing frequencies yields slope $\{\text{python}\} \text{ f"bkl_ll_slope:.3f}"$ (BKL theory: 0.5) with $R^2 = \{\text{python}\} \text{ f"bkl_ll_r_sq:.4f}"$. The near-perfect fit ($R^2 > 0.99$) confirms that variance grows exactly as $1/N$, the hallmark of independent per-step errors that wash out with hedging frequency. Structural errors would manifest as a floor that does not shrink with N .

Floating-point precision. The basis-point conversion routines round to the nearest integer at ≤ 0.5 bp per call. Over $\{\text{python}\} \text{ MC_N_STEPS}$ rebalancing steps with approximately two conversions per step, independent rounding accumulates to a standard deviation of about $\{\text{python}\} \text{ f"rounding_std_dollars :.0f}"$ ($\{\text{python}\} \text{ f"rounding_frac_of_ref:.2f}"\%$ of the BS price), which is $\{\text{python}\} \text{ f"rounding_frac_of_std:.1f}"\%$ of the empirical hedging-cost standard deviation. Basis-point rounding contributes negligibly to observed variance.

8.8 Hull Table 19.2 / 19.3 deterministic cross-check

The Hull (2014) textbook tabulates a deterministic 20-step weekly hedge of a written ATM call ($S_0 = 49$, $K = 50$, $r = 5\%$, $\sigma = 20\%$, $T = 20$ weeks) under two realized price paths: Table 19.2 ends ITM at $S_T = 57.25$; Table 19.3 ends OTM at $S_T = 48.12$. Hull's printed cost-of-hedging figures (\$263{,}300 and \$256{,}600 respectively) incorporate textbook rounding conventions (delta to three decimal places, share counts to the nearest hundred, dollar amounts to the nearest hundred) that introduce noise of several thousand dollars on a ~\$250K hedging cost. The cleaner test is to recompute the hedge in Python using the same basis-point integer arithmetic the Python runtime bridge uses, and compare the two NAV paths point by point.

The two implementations agree to $\{\text{python}\} \text{ f"hull_192_dev_bp:.2f}"$ bp on the ITM path and $\{\text{python}\} \text{ f"hull_193_dev_bp:.2f}"$ bp on the OTM path. The remaining discrepancy is cumulative basis-point rounding through the FFI boundary; the noise floor for end-to-end cross-checks of this kind is about 1 bp per 20-step hedge. Hull's printed cost-of-hedging figures sit roughly 3 to 4% above both implementations; the difference is fully explained by the textbook's

three-decimal-place delta rounding and hundred-share rounding rules. All step audit records pass for both paths.

8.9 QuantLib pricing cross-check

QuantLib is an open-source C++ derivatives library widely used in quantitative finance. Comparing the basis-point integer Black-Scholes pricer to QuantLib's analytic pricer at matched inputs isolates the pricing layer from the position ledger: if the two pricers agree, the cost figures elsewhere in the paper cannot be attributed to a pricing-implementation bug. Practitioners reviewing this work will recognize the comparison as the standard vendor cross-check used by validation desks for new derivatives systems.

Across {python} ql_n_scenarios parameter scenarios the maximum deviation between the basis-point integer pricer and QuantLib's analytic pricer is {python} f"{ql_max_dev:.2f}" bp (?@fig-ql). All {python} ql_n_scenarios scenarios fall within a 3 bp tolerance ({python} "yes" if ql_within_3bp else "no"). Any systematic bias in the integer-arithmetic Black-Scholes implementation large enough to distort the hedging-cost figures elsewhere in the paper would appear as a diagonal offset here; the maximum 2.69 bp deviation rules out pricing bias at the 3 bp threshold used by buy-side validation desks.

8.10 Equity and Covered-Call Demo

The EquityStrategy in the left panel of ?@fig-equity holds 1,000 shares at a fixed quantity with no delta rebalancing. The portfolio value starts at exactly zero (the cost of purchasing the shares is funded by borrowing against the cash account, so the net asset value at inception is zero by selfFinancing) and then tracks $N \times \Delta S$ at each weekly step, with a small interest debit on the borrowed cash. Every step certificate passes, confirming that valueUpdateFormula applies to equity positions without any modification to the Lean kernel: the trade application function receives (assetId, quantity, markPrice) tuples with no instrument-specific fields, so an equity share and an option contract are formally indistinguishable at the ledger layer.

The CoveredCallStrategy in the right panel combines the same equity position with a short ATM call ($K = 50$, $\sigma = 20\%$). The call premium received at inception raises the initial portfolio value above zero. At expiry ($S_T = 57.25 > K$), the call settles in-the-money and settlement_value_formula certifies the cash transfer; the equity leg has no settlement and contributes its final mark-to-market value to the terminal PV. All {python} len(_cc_result.certificates) step certificates pass across the covered-call path, and the terminal PV reflects the capped upside of the covered-call payoff: the equity gain above K is transferred to the call counterparty at settlement. No new Lean theorems were required for either the equity-only or the covered-call strategy: the existing valueUpdateFormula, selfFinancing, and settlement_value_formula compose across instrument types by construction.

8.11 Replication

```
git clone https://github.com/eigenq-xyz/quant-proofs
cd quant-proofs/backtest-proofs

# 1. Verify Lean proofs (Levels 1-2)
make verify-lean

# 2. Run Python tests (Levels 2-4)
make test

# 3. Install research dependencies and generate figures
cd python && uv sync --group research
cd ..
make paper
```

After downloading WRDS data per docs/wrds_data_download.md, re-run `make paper` to regenerate `?@fig-stress` with real market data.

8.12 Metrics snapshot

Numerical results are archived in `backtest-proofs/results/metrics.json`; the file is regenerated on each render.

References

- Annenkov, Danil, and Martin Elsmann, 2018, [Certified compilation of financial contracts](#), *Proceedings of the 20th international symposium on principles and practice of declarative programming (PPDP '18)* (ACM).
- Bahr, Patrick, Jost Berthold, and Martin Elsmann, 2015, [Certified symbolic management of financial multi-party contracts](#), *Proceedings of the twentieth ACM SIGPLAN international conference on functional programming (ICFP 2015)*.
- Bailey, David H., Jonathan M. Borwein, Marcos López de Prado, and Qiji Jim Zhu, 2016, [The probability of backtest overfitting](#), *Journal of Computational Finance* 20, 39–69.
- Bakshi, Gurdip, and Nikunj Kapadia, 2003, [Delta-hedged gains and the negative market volatility risk premium](#), *Review of Financial Studies* 16, 527–566.
- Bertsimas, Dimitris, Leonid Kogan, and Andrew W. Lo, 2000, When is time continuous?, *Journal of Financial Economics* 55, 173–204.
- Board of Governors of the Federal Reserve System, 2011, [Supervisory guidance on model risk management](#), Federal Reserve System / Office of the Comptroller of the Currency.
- Boyle, Phelim P., and David Emanuel, 1980, Discretely adjusted option hedges, *Journal of Financial Economics* 8, 259–282.
- Buehler, Hans, Lukas Gonon, Josef Teichmann, and Ben Wood, 2019, [Deep hedging](#), *Quantitative Finance* 19, 1271–1291.
- Carr, Peter, and Dilip Madan, 1998, Towards a theory of volatility trading, *Volatility: New estimation techniques for pricing derivatives*, 417–427.
- Carr, Peter, and Liuren Wu, 2020, [Option profit and loss attribution and pricing: A new framework](#), *Journal of Finance* 75, 2271–2316.
- Dessalvi, Marco, Massimo Bartoletti, and Alberto Lluch-Lafuente, 2026, A formal approach to AMM fee mechanisms with Lean 4, (arXiv:2602.00101).
- Echenim, Mnacho, Hervé Guiol, and Nicolas Peltier, 2021, [Formalizing the Cox–Ross–Rubinstein pricing of European derivatives in Isabelle/HOL](#), *Journal of Automated Reasoning* 65, 669–714.
- El Karoui, Nicole, Monique Jeanblanc-Picqué, and Steven E. Shreve, 1998, [Robustness of the Black](#)

and Scholes formula, *Mathematical Finance* 8, 93–126.

Figlewski, Stephen, 1989, [Options arbitrage in imperfect markets](#), *Journal of Finance* 44, 1289–1311.

Harvey, Campbell R., Yan Liu, and Heqing Zhu, 2016, [...And the cross-section of expected returns](#), *Review of Financial Studies* 29, 5–68.

Hull, John C., 2014, *Options, Futures, and Other Derivatives*. 9th Global. (Pearson).

Hutchinson, James M., Andrew W. Lo, and Tomaso Poggio, 1994, [A nonparametric approach to pricing and hedging derivative securities via learning networks](#), *Journal of Finance* 49, 851–889.

Jensen, Theis Ingerslev, Bryan T. Kelly, and Lasse Heje Pedersen, 2023, [Is there a replication crisis in finance?](#), *Journal of Finance* 78, 2465–2518.

Kabanov, Yuri M., and Mher M. Safarian, 1997, [On Leland’s strategy of option pricing with transactions costs](#), *Finance and Stochastics* 1, 239–250.

Leland, Hayne E., 1985, [Option pricing and replication with transactions costs](#), *Journal of Finance* 40, 1283–1301.

Löw, Robert, Stanislaus Maier-Paape, and Andreas Platen, 2017, Correctness of backtest engines, *Journal of Investment Strategies*.

Natarajan, Raja, Suneel Sarawat, and Abhishek Kr Singh, 2021, [Verified double sided auctions for financial markets](#), *12th international conference on interactive theorem proving (ITP 2021)*. Leibniz international proceedings in informatics (LIPIcs).

Natarajan, Raja, Suneel Sarawat, and Abhishek Kr Singh, 2024, [Double auctions: Formalization and automated checkers](#), *Journal of Automated Reasoning*.

Passmore, Grant O., Simon Cruanes, Denis Ignatovich, Johannes Kanig, Alain Mebsout, and Seyi Sheridan, 2020, [The Imandra automated reasoning system \(system description\)](#), *Automated reasoning — IJCAR 2020*. Lecture notes in computer science.

Passmore, Grant O., and Denis Ignatovich, 2017, [Formal verification of financial algorithms](#), *Automated deduction — CADE 26*. Lecture notes in computer science.

Peyton Jones, Simon, Jean-Marc Eber, and Julian Seward, 2000, [Composing contracts: An adventure in financial engineering](#), *Proceedings of the fifth ACM SIGPLAN international conference on functional programming (ICFP 2000)*.

Pusceddu, Daniele, and Massimo Bartoletti, 2024, Formalizing automated market makers in the Lean 4 theorem prover, *5th international workshop on formal methods for blockchains (FMBC 2024)*. OASICS.

Shalizi, Cosma Rohilla, 2013, *Advanced Data Analysis from an Elementary Point of View*.

Yin, Dong, Takeshi Miki, Vladislav Lesnichenko, and Vasyl Gural, 2026, [Implementation risk in portfolio backtesting: A previously unquantified source of error](#), (arXiv:2603.20319).